

SQUEEZED ATTENTION: Accelerating Long Context Length LLM Inference

Coleman Hooper^{*1} Sehoon Kim^{*1} Hiva Mohammadzadeh¹ Monishwaran Maheswaran¹
June Paik² Michael W. Mahoney^{1,3,4} Kurt Keutzer¹ Amir Gholami^{1,3}

¹ UC Berkeley ² FuriosaAI ³ ICSI ⁴ LBNL

Abstract

Emerging Large Language Model (LLM) applications require long input prompt in order to perform complex downstream tasks like document analysis and code generation. For these long context length applications, the length of the input prompt poses a significant challenge in terms of inference efficiency since the inference costs increase linearly with sequence length. However, for many of these applications, much of the context in the prompt is fixed across different user inputs, thereby providing the opportunity to perform offline optimizations to process user inputs quickly, as they are received. In this work, we propose SQUEEZED ATTENTION as a mechanism to accelerate LLM applications where a large portion of the input prompt is fixed. To accomplish this, we first leverage K-means clustering offline to group the keys for the fixed context based on semantic similarity and represent each cluster with a single centroid value. During inference, we compare query tokens from the user input with the centroids to predict which of the keys from the fixed context are semantically relevant and need to be loaded during inference. We then compute exact attention using only these important keys from the fixed context. This method maintains model accuracy while significantly reducing bandwidth and computational costs, as exact attention is computed with only a subset of the fixed context tokens. We also extend our method to use a hierarchical centroid lookup to identify important keys, which can reduce the complexity of attention from linear to logarithmic with respect to the fixed context length. To realize our method's efficiency benefits, we implement optimized Triton kernels for centroid comparison and sparse FlashAttention with important keys, achieving more than $4\times$ speedups during both the prefill and generation phases for long-context inference. Furthermore, we have extensively evaluated our method on various long-context benchmarks including LongBench, where it achieves a $3.1\times$ reduction in KV cache budget without accuracy loss. For applications where small accuracy degradation is allowed, we can achieve up to an $8\times$ reduction with less than 0.5 point accuracy gap for the LLaMA-2-7B-32K, LWM-Text-Chat-1M, and Longchat-7B-v1.5-32K models. Our code is available at <https://github.com/SqueezeAILab/SqueezedAttention>.

1 Introduction

Large Language Models (LLMs) have seen rapid advancements in recent years, enabling a range of downstream applications including document Question Answering (QA) and complex analysis over structured and unstructured documents. Performance on these document processing and analysis tasks has benefited from the increased context lengths of newer open-source [26, 40] and closed-source [3, 4, 15] models, as these tasks benefit from incorporating a large amount of input context in order to condition the model to generate particular outputs. However, deployment of LLMs for downstream applications is constrained by inference costs, with LLM inference requiring significant computational resources as well as memory capacity and bandwidth. In particular, long context-length applications have large memory capacity and memory bandwidth requirements due to the size of the KV cache, which increases linearly with respect to sequence length [39, 17, 24].

For many applications such as in-context learning, document QA, and code generation, over a series of prompts a large portion of the input context is fixed. For example, the input context may contain system instructions, documentation, source code, as well as particular few-shot in-context examples for the target task. This fixed context, which

^{*}Equal contribution




Figure 1: A high-level visualization of our hierarchical clustering approach. We identify important keys for the current query by first identifying which coarse-grained clusters are relevant (Level 1). We then refine this prediction using finer-grained clustering (Level 2). Finally, we identify the important keys for the current query and only compute exact attention with these keys.

is reused in successive prompts, is extremely beneficial for tailoring the model to the target application; however, increasing the size of the fixed context poses a significant challenge for inference efficiency. Throughout this work, we will refer to this portion of the prompt as the “fixed context,” and we will refer to the portion that corresponds to the user requests that come in online as the “user input.” The user input is appended after the fixed context and provided to the model. For many long-context applications, the fixed context portion of the prompt is much longer than the user input portion of the prompt, and the attention computation for this fixed context portion typically dominates the inference runtime. In this work, our aim is to take advantage of the fact that this context is fixed and available prior to inference. This allows us to optimize the attention to this fixed context when processing incoming user inputs and generating outputs.

To this end, we propose SQUEEZED ATTENTION as a method to accelerate fixed context applications by accelerating the attention computation. Our method, illustrated in Figure 1, accelerates inference by quickly identifying which keys in the fixed context are important for a given query token. This approach involves a two-stage algorithm. In the first stage, we cluster the keys in the fixed context offline based on their semantic similarity and then represent keys from the same cluster using a single representative “key centroid”. This offline clustering enables fast retrieval of important keys (i.e., those most semantically relevant to incoming query tokens) during inference. In the second stage, when the user input is received, we retrieve the important keys by first comparing the query tokens with the key centroids, rather than the entire set of keys, in order to identify the important key clusters. Once the important clusters are identified, we retrieve their associated keys and compute exact attention only with those high-scoring keys. Our method can be further extended to a *hierarchical* clustering and retrieval scheme, as shown in Figure 1, efficiently narrowing the search space by first leveraging coarser-grained clusters and then refining the search using fine-grained clusters. As we will later discuss in Section 3.4, this approach can reduce the memory and compute complexity of lookups to logarithmic complexity with respect to the fixed context length.

In contrast to existing solutions [46, 24, 14] that identify less important tokens once and discard them throughout the entire generation, our method *dynamically* identifies and retrieves only the information that is *semantically* relevant to each generation step. This allows our method to preserve generation quality while reducing the number of KV cache entries loaded from memory by up to 8 times (including loading key centroids), as highlighted in Section 6. By optimizing memory bandwidth as well as computational costs, SQUEEZED ATTENTION effectively reduces overheads for both generation and prefill during long-context inference. Specifically, our work makes the following contributions toward accelerating long context length applications:

- **Semantic-based Key Clustering and Retrieval:** To cluster non-consecutive keys by their semantic similarity, we perform K-means clustering offline, representing all keys within each cluster with a single “key centroid” value (Section 3.1). This enables efficient retrieval during inference, allowing us to identify the keys most semantically relevant to the query tokens by comparing the query against key clusters instead of the entire key set (Section 3.2). Since the number of key centroids is significantly smaller than the total number of keys, the memory overhead

remains minimal. Then, the exact attention scores are computed only with the retrieved keys. We additionally propose the hierarchical version of our method, which can reduce the memory and computational complexity of per token generation from linear to logarithmic with respect to the context length.

- **System Implementation:** To realize the efficiency benefits of our approach, we design efficient Triton kernels for performing the centroid comparison (Section 4.1) and computing sparse FlashAttention with only the important keys (Section 4.2). Combined together, our method results in $4.3\times$ and $4.2\times$ speedups during the prefill and decode phases when running inference with long fixed context. For applications where small accuracy degradation is allowed, we can achieve up to an $8\times$ reduction with less than 0.5 point accuracy gap for the LLaMA-2-7B-32K, LWM-Text-Chat-1M, and Longchat-7B-v1.5-32K models (Section 6.3).
- **Benchmark:** There is currently few long context QA benchmark datasets which represent applications where the user is asking different questions against a document/knowledge source. To address this, we introduce PreFixQA, a document QA benchmark which contains a selection of arXiv documents, each with many synthetic user input questions and answer pairs against a fixed knowledge source. This benchmark facilitates research into fixed context methods by allowing us to evaluate various user inputs for each document (Section 5).
- **Evaluation:** We extensively evaluate our method on different long-context benchmarks including LongBench [5], RULER [18], and PreFixQA. Particularly, on LongBench, our method preserves the full KV cache accuracy with $3.1\times$ KV budget reduction. For applications that can tolerate a small accuracy degradation, SQUEEZED ATTENTION achieves up to $8\times$ KV budget reduction with less than 0.5 point accuracy drop (Section 6).

2 Related Work

2.1 Long-Context LLMs

With the growing popularity of long-context applications, there has been a continuous development of LLMs that can support context lengths exceeding 100K, and even up to 1M tokens. This includes proprietary models such as GPT-4-Turbo [3], Claude-2 [4] and Gemini 1.5 [15], which support context lengths of up to 128K, 200K, and 1M tokens, respectively. On the open-source front, several efforts have been made to extend the context lengths beyond the length on which the original models were trained [23, 7]. A notable work is Large World Model (LWM) [26], which has demonstrated extending the context length of Llama 2 [41] to 1M tokens. However, as context lengths increase, the KV cache often becomes a critical bottleneck, significantly impacting memory usage and latency during LLM inference [39, 17]. Therefore, KV cache compression methods have emerged as a critical concern for enabling efficient inference when using long-context models.

2.2 KV Cache Compression for Long-Context Inference

To enable more efficient long-context inference by reducing the KV cache size, several methods have been proposed, including quantization [17, 29, 19, 27], shared KV cache across tokens [32] and layers [6], and token pruning [13]. A notable approach which will be discussed in more detail is KV cache sparsification, which follows a prior line of work in attention sparsification [38, 8, 44]. There are two general directions which have been pursued for KV cache sparsification: KV cache eviction, and sparsely loading the KV cache.

KV Cache Eviction. KV eviction has become a widely used method for compressing the KV cache by identifying and removing less important tokens. Various strategies have been proposed to determine token importance, including attention score contribution [46, 35], persistent attention patterns during generation [28], token entropy [43], and additional heuristic-based policies [14].

However, in use cases where long context prompts are followed by varying questions, the importance of the KV cache for the context should be decided on the basis of its relevance to the subsequent question. To address this, SnapKV [24] proposes selecting KV cache entries solely based on the attention scores of the most recent prompt tokens to the rest of the input prompt. However, since the important tokens in the input prompt are determined once and remain fixed throughout the generation process, it cannot adapt to changing token importance during generation or in response to subsequent user inputs. InfiniPot [21] extends this idea by iteratively compressing the context based on its relevance to predefined task-specific prompts that resemble potential input questions. Nevertheless, selecting important tokens offline using proxy prompts may not accurately reflect future queries.

shortage



Likewise, eviction-based approaches discard tokens and retain the remaining ones throughout generation, potentially overlooking the fact that discarded tokens could become important later in the process. SQUEEZED ATTENTION, on the other hand, bypasses the need for a full KV cache lookup by clustering the KV cache and retrieving only the most relevant clusters through an efficient centroid lookup. This approach is lightweight enough to be applied at every generation step, thereby ensuring relevant context is retrieved for every query token.

Sparse KV Cache Loading. One previous direction that has been explored aims to store the full KV cache, but only load in the relevant keys and values dynamically during inference. QUEST [39] clusters consecutive KV cache entries and dynamically retrieves the most relevant clusters based on their relevance to each query token during generation.

Another line of relevant work here is application of fast kernel summation methods [16, 42, 22, 31, 30] and in particular variants of Fast Multipole Method (FMM) [9] which were originally proposed to accelerate N-body simulations. In the context of Transformers, recent work of [20] utilizes FMM to cluster consecutive past tokens and assign coarser-grained clusters to older tokens, reducing the memory overhead of storing the entire past tokens. However, this approach, as well as QUEST [39], rely on *physical* proximity for clustering, whereas in Natural Language Applications clustering should instead be based on *semantic* proximity, as tokens that are physically far apart can be semantically similar. This is because tokens that are far apart can be semantically relevant and vice versa. SQUEEZED ATTENTION addresses this by clustering tokens based on their embedding similarity, ensuring that semantically relevant tokens are retrieved for future generations.

Another prior line of work aims to leverage vector search methods for only loading important keys and values. PQCache [45] applied product quantization-based vector search to identify important keys. RetrievalAttention [25] uses a K-Nearest Neighbors-based vector search approach, which offloads dynamic retrieval of important keys and values to the CPU. However, these prior approaches are restricted to the generation stage and do not accelerate prefill, which is critical to reducing time-to-first-token (TTFT) latencies.

In contrast with prior works which leverage vector search methods, SQUEEZED ATTENTION uses a fast centroid lookup to enable accurate retrieval of relevant contexts on the GPU without requiring offloading operations to CPUs, as in [25]. Our approach is also able to accelerate both prefill and generation. Furthermore, our method allows for loading more or fewer keys from different heads, depending on the number of important keys for each head. This approach enables us to achieve higher accuracy while aggressively reducing the number of KV entries.

3 Algorithm

We design our method to preprocess the fixed context offline, so that at inference time we can quickly determine which information is important and only load this information. In Section 3.1, we discuss how we cluster the fixed context keys offline based on their semantic similarity and then determine a representative centroid for each cluster. In Section 3.2, we propose a method to identify which clusters are the most important to retrieve for precise attention computation based on input queries during inference. Finally, in Section 3.3, we extend our algorithm to include multiple levels of centroids in order to accelerate the search for important key tokens, thereby improving the scalability of our method for longer context lengths.

3.1 Offline: Clustering Keys

The first step in our method is to preprocess the fixed context keys offline, as outlined in Figure 2. We take the fixed context keys and cluster them based on cosine similarity. Specifically, we use K-means clustering with normalized key vectors to group similar keys together. We then compute a centroid for each cluster by taking the mean of all of the vectors in the cluster. This cluster centroid can be used as a representative key token for all tokens in that cluster; by comparing incoming queries with this centroid, we can determine whether the tokens in that cluster are important without necessarily comparing them with individual keys.

Note that this *semantic-based* clustering approach groups together non-consecutive key tokens, which could make it more challenging to efficiently load the keys from memory. However, the size of each KV cache token for a single head in modern LLMs is typically larger than 256 bytes in bf16 (as the head dimensions are typically greater than 128) [40, 41], which is sufficiently large to efficiently utilize memory bandwidth. Therefore, we are still able to execute memory operations efficiently when sparsely loading in the non-consecutive keys and associated values from memory.




Figure 2: Diagram outlining our approach for performing clustering offline with the fixed context. Refer to Section 3.1 for 1-level clustering and Section 3.3 for hierarchical clustering. We apply K-means clustering to group semantically similar key tokens, assigning a single centroid to represent each cluster. In the hierarchical approach (Section 3.3, demonstrating a 2-level hierarchy for clarity), these centroids form the Level 2 centroids, which are then clustered into coarser-grained Level 1 centroids by repeating the same procedure.

3.2 Online: Query-Aware Key Retrieval

Ideally, we would load only the keys whose attention scores are high. However, which keys will be high-scoring for a given query cannot be known ahead of time, without doing a full pass over the keys. In our approach, we use the centroid cluster to approximately measure the “average” attention score of the keys within a cluster, thereby allowing us to identify important keys without loading them all.

By organizing keys into clusters, each represented by a single centroid, we can accelerate inference for incoming user inputs, as illustrated in Figure 3. We first compare the input query tokens with each of the key centroids to assess which key tokens are likely to have high attention scores. To estimate the importance of cluster i for query token q , we compute the attention estimate for that cluster as:

$$S_i = \frac{\exp(qC_i^\top)}{\sum_j N_j \cdot \exp(qC_j^\top)}, \quad (1)$$

where N_j is the number of keys in cluster j and C_j is the cluster centroid for cluster j . This allows us to assess the average importance of the tokens in cluster i . If the average importance of a cluster is above a desired threshold, we load in the keys for that cluster and perform exact attention computation; otherwise, we avoid loading and performing computation with these keys.

Using the Softmax estimate S_i , instead of qC_i^\top , as an importance metric for each cluster provides an easy method to control the number of important keys retrieved from each attention head. As outlined in Appendix A.2, some attention heads have a more balanced distribution of attention scores, resulting in a larger number of important keys, while others have a more skewed distribution, indicating only a few important keys. Ideally, we want to retrieve more keys from heads with a larger number of important keys. Since Softmax values are normalized to sum to 1, we can apply a single *global* threshold across all layers and attention heads to achieve this. This allows us to automatically retrieve more keys from heads with balanced attention score distributions, where more S_i values exceed the threshold; and fewer keys from heads with skewed distributions, where fewer S_i values exceed the threshold. This approach eliminates the need for manually configuring the number of keys to retrieve for each head. Once we choose the threshold to achieve the desired sparsity level, it is kept throughout the prefill and generation stages.

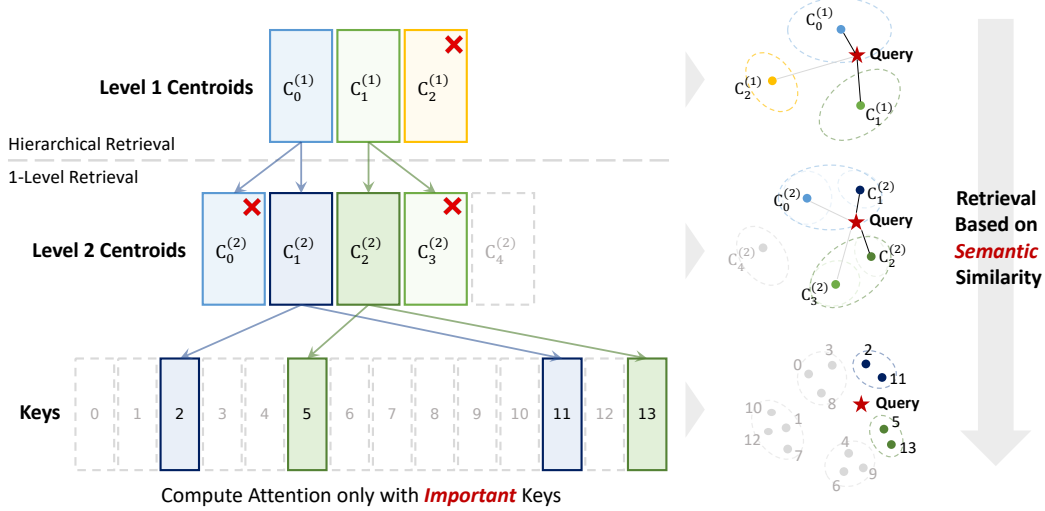


Figure 3: Diagram outlining how our method operates during inference to retrieve the most relevant keys when a new input query is received. Refer to Section 3.2 for 1-level retrieval and Section 3.3 for hierarchical retrieval. For 1-level retrieval, the query token is first compared against the representative centroid of each cluster to identify the most relevant clusters. Exact attention is then computed only for the keys within these retrieved clusters, rather than across the entire fixed context. In our hierarchical retrieval approach (Section 3.3, demonstrating a 2-level hierarchy for clarity), we first compare the query with coarse-grained Level 1 centroids, and then only compare with a subset of the promising fine-grained Level 2 centroids in order to identify the important keys.

3.3 Hierarchical Centroid Lookup

The centroid lookup approach outlined in Sections 3.1 and 3.2 allows for quickly determining which keys are likely to be important in the attention computation, and then only computing attention with these keys. As long as we use fine-grained centroids, we can have sufficient resolution to identify which keys will be important and we can retain accuracy. However, it is desirable to keep a smaller number of centroids, since a larger number of centroids leads to an increased cost for centroid lookup.

In order to attain the accuracy improvements of fine-grained centroid lookup while retaining the efficiency benefits of using coarse-grained centroids, we leverage a *hierarchical* centroid lookup process. Figure 2 demonstrates our (two-level) hierarchical approach during offline preprocessing of the fixed context keys. Initially, using the same approach as in Section 3.1, we cluster the keys into a larger number of centroids, referred to as Level 2 centroids. Then, we perform an additional K-means clustering on these Level 2 centroids to produce a smaller number of coarse-grained centroids, referred to as Level 1 centroids.

During inference, we perform a hierarchical centroid lookup as outlined in Figure 3. We first compare incoming queries with the coarse-grained Level 1 centroids to quickly prune out unnecessary keys. This initial lookup narrows down the search, allowing us to focus on comparing the queries with the fine-grained Level 2 centroids that are likely to be high-scoring. Specifically, we first compare the input query token q with each of the coarse-grained key centroids $C_i^{(1)}$ to assess which key tokens are likely to have high attention scores:

$$S_i^{(1)} = \frac{\exp(qC_i^{(1)\top})}{\sum_j N_j^{(1)} \cdot \exp(qC_j^{(1)\top})}. \quad (2)$$

We then apply a threshold T_1 to rule out low-scoring clusters at the coarse-grained level. This filtering allows us to avoid performing any comparisons for the fine-grained Level 2 centroids for the tokens which are unlikely to be high-scoring. For the remaining Level 1 centroids, we expand them into their corresponding finer-grained Level 2

Table 1: Theoretical memory and compute complexity of the baseline (standard autoregressive generation), 1-level retrieval, and hierarchical retrieval for a single generation iteration. Here, L represents the context length, c is the number of clusters in the 1-level retrieval approach, and $k \ll L$ is the number of keys remaining after retrieval. In the hierarchical retrieval approach, $c' \ll c < L$ denotes the number of clusters at each hierarchical level. Note that c for the 1-level retrieval cannot be reduced significantly with respect to L , while c' for the hierarchical retrieval can.

Method	Memory / Compute Complexity
Baseline	$O(L)$
1-Level Retrieval	$O(c + k)$ where $k \ll L$
Hierarchical Retrieval	$O(c' \log L + k)$ where $c', k \ll L$

centroids, $C_l^{(2)}$, which are then compared with the input query token q to assess their relevance:

$$S_l^{(2)} = \frac{\exp(q C_l^{(2)\top})}{\sum_m N_m^{(2)} \cdot \exp(q C_m^{(2)\top})}. \quad (3)$$

Since we are only considering the remaining Level 2 centroids, the denominator is also calculated based on these selected centroids. We then compare $S_l^{(2)}$ with threshold T_2 to decide which keys should be used for exact attention computation. With this hierarchical approach, we can reduce the cost of finer-grained centroid lookup while maintaining its accuracy. Although we describe a 2-level process here for clarity, this method can be extended to multiple levels of hierarchy.

3.4 Complexity Analysis

Let L denote the context length, which can be substantially large in long-prompt applications. In the baseline approach (i.e., standard autoregressive generation), each generative step requires comparing a query token with the entire set of keys in the prompt, resulting in $O(L)$ memory and compute operations per iteration (i.e. per token generation). If we apply 1-level retrieval, however, we can instead use c centroids to identify the relevant key clusters and then compute attention using only $k \ll L$ retrieved keys. This reduces the memory and compute complexity to $O(c + k)$ per iteration. One limitation of the 1-level retrieval approach is that it can be challenging to significantly reduce c (the number of centroids), as it would require clustering a large number of keys into each cluster. This may result in either pruning keys too aggressively or retrieving irrelevant keys grouped together in the same cluster.

In contrast, hierarchical centroid retrieval allows for a more efficient reduction in centroids at each level of the hierarchy by enabling gradual pruning of keys. Suppose we use $c' \ll c < L$ clusters at each hierarchical level and retrieve only a fraction, $0 < p < 1$, of these clusters at each stage. In this setup, we require $O(\log L)$ hierarchical levels to reduce the keys to the desired final count, k . Therefore, the memory and compute complexity for each generation iteration becomes $O(c' \log L + k)$, reducing the complexity from *linear* to *logarithmic* with respect to the context length.

4 System Implementation

In order to realize the efficiency benefits of our method, we designed Triton [33] kernels to compute each stage of our pipeline efficiently online during inference. The first stage (Section 4.1) computes the centroid lookup in order to determine the tokens for which we must compute attention exactly, based on the algorithms discussed in Section 3. The second stage (Section 4.2) leverages this information to (i) only load in the important keys and to (ii) only perform the attention computation with the important keys, thereby saving both compute and memory bandwidth. For long-context applications with long, fixed context, the KV cache for the fixed context can be cached separately from the KV cache that is dynamically generated by incoming user inputs.

4.1 Centroid Lookup

The first stage of our kernel implementation compares query tokens with the centroids for the fixed context keys. These query tokens may include multiple tokens from an incoming user input during the prefill stage or a single token during

the generation stage. The kernel follows similar parallelization strategies to FlashAttention-2 [10], where we split across different attention heads and along the query sequence length dimension. We first load a block of query tokens and iterate over the entire key centroids in order to find the most important key centroids according to Equation 1. At a high level, the kernel performs an initial pass over the key centroids to compute the denominator in Equation 1 based on the query-key centroid dot product. Then, it takes a second pass over the centroids to compute S_i as in Equation 1, using the denominator results from the first pass. Finally, we compare S_i with a target threshold T , and we only load the keys in cluster i if $S_i > T$. Details of how this process is applied during the prefill and generation stages will be discussed in the following subsections.

Prefill Stage. During prefill, where multiple query tokens are available, we split the workload along the query sequence length dimension as in FlashAttention-2 [10] to attain additional parallelism. Since this process produces individual S_i values for each query token, we compute their average to obtain \bar{S}_i , i.e., the *averaged* importance score for each key cluster across all query tokens. We then check whether $\bar{S}_i > T$ to determine whether to load the keys in the corresponding cluster. Since S_i is an estimate of the Softmax value, which is normalized to sum to 1, averaging across query tokens provides a simple way to calculate their combined importance score.

Generation Stage. During generation, achieving parallelism is more challenging, as we cannot leverage parallelism across the dimension of the length of the query sequence. This is particularly problematic when dealing with small batch sizes, as in that case the only parallelism we can leverage is across different heads. To accelerate centroid lookup during generation, we additionally compute and store $\exp(qC_i^\top)$ for each cluster during the first pass over the key centroids while we are computing the denominator $D = \sum_j N_j \cdot \exp(qC_j^\top)$. Then, in the second pass, we load the precomputed $\exp(qC_i^\top)$ values and compare them against DT to determine the importance of each cluster, without the need to explicitly compute S_i . This second pass can be parallelized across the cluster dimension for fast comparison. As highlighted in Appendix B, these optimizations are crucial for performing the centroid lookup during generation without substantial latency overhead.

4.2 Sparse Attention with Retrieved Keys

Once the important keys are identified through our centroid lookup, the second stage of our system implementation leverages a sparse FlashAttention kernel to attain speedups during *both* prefill and generation stages. This stage also uses similar parallelization strategies as FlashAttention-2 by splitting work across heads and along the sequence length dimension [10]. Our kernel implementation builds on top of prior work on Triton implementations for FlashAttention-2 [34] and for dynamic sparse FlashAttention [36]. The kernel first loads in query vectors, and then iterates over a tensor of key indices that need to be selectively loaded. These indices are then used to load the corresponding keys from memory and compute exact attention.

An additional challenge when computing attention to the fixed context is the imbalanced distribution of important key tokens across different heads, which is highlighted in Figure A.2 in Appendix A.2. When using the default parallelization strategy in FlashAttention-2, if one head contains more important keys than the other heads, it will have significantly longer runtime, hindering speedups. In order to obtain latency benefits in these scenarios, we split keys and values along the sequence-length dimension as in Flash-Decoding [11], based on a fixed number of desired keys and values to be computed for a single Streaming Multiprocessor (SM). This means that if there are more keys and values that need to be computed for a particular head (due to unbalanced sparsity for different heads), the work for this head will be parallelized across a greater number of SMs in the GPU. The kernel is designed in two phases, as in Flash-Decoding. The first phase computes the partial attention outputs for each block of valid keys and values. The second stage merges the partial attention outputs, while correcting the outputs using the partial Softmax denominators and max values.

5 Benchmark for Fixed Context Processing

Despite the growing demand for long-context applications where a fixed document is used to answer multiple user requests (e.g., code generation or long-document QA), there is currently no benchmark designed to test this scenario. Recent long context benchmarks, such as LongBench [5] and RULER [18], only pair each long-context input with a single question. These benchmarks do not evaluate the handling of multiple queries on the same document, which presents a challenge for developing fixed context optimization methods, as it leads to a longer iteration cycle. This




Figure 4: Visualization of samples from PreFixQA. The dataset consists of a set of arXiv documents, along with a series of questions for each sample document.

is because the offline preprocessing step would need to be performed for every single sample instead of once per fixed context, which could otherwise be reused across multiple samples where the fixed context is shared. To bridge this gap, we introduce a new benchmark called PreFixQA (Prefix-Fixed QA) which evaluates the ability of LLMs to manage multiple queries on a single long-context document. Our dataset curation pipeline consists of two phases: (i) collecting long documents from arXiv papers; and (ii) generating question-answer pairs based on each document, while ensuring correctness and consistency, which will be discussed in detail below.

Long Documents Collection. PreFixQA is a long-document QA benchmark designed for one-document-multi-user-input scenarios. To collect high-quality long documents, we have sampled 47 papers from arXiv, each ranging from 17,000 to 200,000 characters, with an average length of 20 pages each after deleting references and appendices. To evaluate LLMs’ capability to understand diverse types of content, we have selected papers from various fields, including computer science, electrical engineering, biology, machine learning, economics, and finance. To prevent training set contamination, all papers were sourced from 2024 using the arXiv API [1].

Question and Answer Generation. To generate multiple questions per each document, we have implemented a multi-step generation and filtering procedure inspired by the Llama-3 training data generation approach [12]. To ensure that questions cover different sections of the document and avoid redundant questions by focusing too heavily on one part, we divide each document into multiple chunks. Each chunk is then provided to GPT-4-Turbo [3] to generate potential questions that can be answered in 1-2 words. We generate questions with short answers to enable more accurate evaluation through string comparison.

However, a single pass of question generation often results in low-quality question-answer pairs due to incorrectness or inconsistency of the answers. To avoid this, we introduce an additional filtering process to ensure the correctness and consistency of each question-answer pair. In this step, each question (along with the specific chunk) is provided to GPT-4-Turbo five separate times to produce potential answers. We then filter out questions with inconsistent answers, which typically arise from ambiguity in the question or context. Furthermore, GPT-4-Turbo is used as a judge to score the correctness of each answer, based on the full document, on a scale of zero to ten. Questions where at least three out of five answers score above eight are kept; otherwise, they are discarded. For the retained questions, the highest-rated answer is selected and kept in the dataset. The prompts used for dataset generation and filtering are provided in Appendix D. This process has yielded 1,127 high-quality question-answer pairs (24 on average per document) for our benchmark, ensuring a diverse and challenging benchmark for long-document question-answering.

6 Results

6.1 Experimental Details

We evaluate our method on a range of downstream long context length tasks. We leverage the LongBench [5] and RULER [18] benchmark suites as well as PreFixQA for our evaluation. Across all tasks, when identifying the “fixed context” portion of the input, we isolate the context before the user input using the prompt template for each task, and then we apply our approach to this fixed portion of the prompt. For our single-level experiments, we set the number of cluster centroids to be 5% of the fixed context length, and for our hierarchical experiments we set the number of Level 1 centroids and Level 2 centroids to be 1% and 5% of the fixed context length, respectively. For our hierarchical experiments, we set the Level 1 threshold such that 50% of the keys would be ruled out before performing

Table 2: LongBench evaluation results with SQUEEZED ATTENTION. We report results across different LongBench tasks for the Llama-2-7B-32K, LWM-Text-Chat-1M, and Longchat-7b-v1.5-32K models, using both our single-level (“Squeeze”) and hierarchical (“H-Squeeze”) lookup approaches. We also report baseline comparisons with QUEST [39], demonstrating how our use of semantic similarity when clustering keys outperforms grouping keys sequentially. We report the average score across LongBench tasks, as well as the average without SamSum (“SSum”) for comparison against QUEST, whose evaluation framework does not support SamSum. We also include the KV budget (“Budget”), which gives the expected percentage of the KV cache that needs to be loaded in during inference (including extra memory movement for cluster centroids). Additional experimental details are provided in Appendix C.

Config	Budget	Single-Document QA			Multi-Document QA			Summarization			Few-shot Learning			Code		Avg.	
		NQA	Qasper	MFQA	Hotpot	2Wiki	Mustique	GRep	QMSum	MNews	TREC	TQA	SSum	RBench	LCC	w/o SSum	All
LLaMA-2-7B-32K																	
All KV	1	17.91	11.12	33.87	12.45	11.95	6.54	29.37	16.93	21.58	71.50	87.96	43.87	61.45	59.14	33.98	34.69
Squeeze-70%	0.325	18.55	11.78	34.33	12.31	12.31	6.26	29.50	16.90	20.76	69.00	87.96	43.90	61.29	59.53	33.88	34.60
QUEST	0.215	17.01	9.89	32.10	11.94	11.41	6.27	28.90	17.65	22.14	68.00	86.43	-	62.53	59.39	33.36	-
Squeeze-80%	0.225	19.03	12.11	32.77	12.51	11.53	6.66	28.82	17.19	20.70	69.00	87.46	44.42	61.26	59.78	33.76	34.52
QUEST	0.168	20.42	9.72	29.46	11.45	9.75	5.46	27.06	17.20	21.83	68.50	86.36	-	61.93	59.38	32.96	-
Squeeze-90%	0.125	18.15	14.39	32.38	11.84	11.70	6.45	29.06	16.93	21.66	70.00	87.43	45.15	58.79	59.37	33.70	34.52
H-Squeeze-90%	0.112	17.41	14.23	32.71	11.99	11.38	6.68	29.14	16.97	20.41	68.00	87.37	44.85	58.94	59.61	33.45	34.26
LWM-Text-Chat-1M																	
All KV	1	16.27	24.36	42.00	21.63	16.70	9.10	27.57	24.71	24.48	70.50	61.70	39.59	41.77	40.72	32.42	32.94
Squeeze-70%	0.325	16.54	24.71	42.24	21.66	15.88	9.08	27.28	24.77	24.60	70.50	60.93	39.75	41.06	40.76	32.31	32.84
QUEST	0.215	15.24	24.57	40.68	21.57	17.02	7.93	27.29	24.86	24.45	67.00	62.14	-	45.53	43.48	32.44	-
Squeeze-80%	0.225	16.66	24.70	41.88	21.10	15.91	9.13	27.00	24.68	24.23	70.00	60.81	39.37	42.07	41.89	32.31	32.82
QUEST	0.168	15.37	23.33	41.45	20.26	17.39	7.85	25.88	25.06	24.43	65.00	62.54	-	46.20	43.06	32.14	-
Squeeze-90%	0.125	16.97	24.96	41.14	20.70	16.40	9.24	27.00	24.59	23.51	71.50	59.37	39.87	44.78	43.80	32.61	33.13
H-Squeeze-90%	0.118	16.69	24.79	40.38	20.78	16.21	8.91	25.02	24.77	22.34	70.50	58.23	39.40	44.31	43.34	32.02	32.55
LongChat-7B-v1.5-32K																	
All KV	1	20.82	28.95	43.06	32.79	24.18	14.09	30.67	22.83	26.09	66.50	83.45	41.25	53.20	56.64	38.71	38.89
Squeeze-70%	0.325	20.93	29.18	43.00	33.02	23.61	14.55	31.13	22.93	26.25	66.50	83.60	40.90	54.64	56.93	38.94	39.08
QUEST	0.215	19.33	31.51	41.65	31.79	23.25	12.58	31.09	22.84	26.87	67.50	84.33	-	53.57	55.37	38.59	-
Squeeze-80%	0.225	20.57	29.64	42.80	33.06	23.63	15.27	31.31	23.21	26.17	65.50	83.87	41.28	52.83	57.17	38.85	39.02
QUEST	0.168	18.03	30.21	37.83	31.78	21.03	11.21	30.52	22.84	26.47	63.50	84.71	-	51.50	55.82	37.34	-
Squeeze-90%	0.125	18.60	29.86	42.21	35.71	23.12	14.31	31.61	22.79	26.17	65.50	78.85	41.22	51.57	56.95	38.25	38.46
H-Squeeze-90%	0.122	18.86	30.51	42.25	35.42	20.88	13.85	30.85	22.84	25.71	65.50	78.50	40.96	51.89	57.20	38.02	38.23

the fine-grained Level 2 lookup. For measuring accuracy with PreFixQA, we use F1 Score to calculate the similarity score between the outputs and the ground truth (as is used in LongBench for single-document QA tasks [5]). We use 32K as the maximum context length throughout our evaluation, and we truncate longer inputs for both LongBench and PreFixQA.

6.2 Accuracy Evaluation Results

6.2.1 LongBench

In order to validate our method, we perform evaluation on long-context datasets from LongBench [5], a comprehensive multi-task benchmark for long-context understanding. LongBench contains a broad set of tasks that cover a range of key long-context applications including Single and Multi-Document QA, Few-shot Learning, Summarization, and Code Completion. We evaluate our approach on the non-synthetic English language tasks in Longbench. Table 2 provides evaluation of our method on LongBench for the LLaMA-2-7B-32K [2], LWM-Text-Chat-1M [26], and Longchat-7B-v1.5-32K models [23]. We also provide baseline comparisons with QUEST [39]. To ensure a fair comparison, we set the token budget for their method dynamically for each input sample to match our approach. Additional details for the baseline configuration are provided in Appendix C.

The results show that our method provides similar accuracy as the full KV cache baseline on long context-length

Table 3: RULER evaluation results with SQUEEZED ATTENTION. We report results across different RULER tasks for the LWM-Text-Chat-1M model using 32K context length for evaluation. We also report the KV budget (“Budget”), which gives the expected percentage of the KV cache that needs to be loaded in during inference for each configuration. Our results show that our method is able to retain the accuracy of the baseline model, even with aggressive sparsity settings.

Config	Budget	Niah1	Niah2	Niah3	MKey1	MKey2	MKey3	MValue	MQuery	VT	CWE	FWE	QA1	QA2	Average
All KV	1	100.0	100.0	99.4	100.0	99.6	96.4	45.6	35.3	58.4	9.7	66.7	63.2	43.6	70.6
Squeeze-70%	0.325	100.0	100.0	99.2	100.0	99.6	93.2	43.8	37.4	57.7	8.4	65.8	60.2	42.0	69.8
Squeeze-90%	0.125	100.0	100.0	98.6	100.0	99.4	81.8	42.1	36.0	51.4	9.7	65.1	57.8	41.6	68.0

Table 4: PreFixQA evaluation results with SQUEEZED ATTENTION. We report results using our single-level approach (“Squeeze”) and with our hierarchical lookup-based approach (“H-Squeeze”). LLaMA-2, LWM, and LongChat are LLaMA-2-7B-32K, LWM-Text-Chat-1M, and LongChat-7B-v1.5-32K, respectively.

Config	LLaMA-2	LWM	LongChat
All KV	43.47	14.92	24.13
Squeeze-70%	42.14	14.45	23.79
Squeeze-90%	36.95	14.25	23.94
H-Squeeze-90%	37.05	14.12	23.71

tasks, while offering significant efficiency improvement in terms of reduction in KV cache loading and attention computation. Across all three models, our method maintains full KV cache accuracy with less than 0.11 point degradation at 70% sparsity, reducing the KV budget by $3.1\times$. Even at a more aggressive 90% sparsity, which reduces the KV budget by $8\times$, our method only introduces a small accuracy degradation of within 0.5 points. Note that our method’s accuracy also matches an idealized baseline, where full attention is computed with all keys before retaining only the highest-scoring ones, as further discussed in Appendix E. This demonstrates that our method can effectively identify and retrieve only the most relevant keys (i.e., those that yield high attention scores) without loading the entire keys.

Furthermore, our method outperforms the QUEST baseline, with a pronounced accuracy gap of up to ~ 1 point for more aggressive sparsity settings. This highlights the advantage of semantic-based clustering, which allows for more aggressive KV cache sparsity by loading only the important keys without losing critical information. Additionally, we include results for our hierarchical lookup approach with 90% sparsity, demonstrating how our hierarchical method has lower overhead from the centroid lookup with minor accuracy loss relative to performing a single-level lookup.

6.2.2 RULER

We also present an evaluation for our method on the RULER benchmark [18] in Table 3, which is a synthetic benchmark designed to provide a comprehensive evaluation of long-context capabilities of language models. The benchmark consists of 13 tasks grouped into four categories: Retrieval, Multi-Hop Tracing, Aggregation, and Question Answering. We use the default RULER configuration with 500 samples to evaluate our method, and we evaluate on RULER using the LWM-Text-Chat-1M model [26]. The results on RULER demonstrate how our method provides similar accuracy as the baseline on long context-length tasks.

6.2.3 PreFixQA

We present evaluation for our method on this dataset in Table 4 using the LLaMA-2-7B-32K, Longchat-7B-v1.5-32K, and LWM-Text-Chat-1M models [2, 23, 26]. The results demonstrate how our method provides similar accuracy as the baseline for fixed context use-cases, while significantly compressing the fixed context that needs to be dynamically loaded during inference.

6.3 Systems Results

6.3.1 Experiment Details

In order to benchmark our kernel implementations for long context length inference, we used sample text from the PG-19 language modeling dataset [37], and applied our clustering method to derive centroids to use when benchmarking.




Figure 5: Kernel implementation latency results for FlashAttention baseline as well as for SQUEEZED ATTENTION with 70%, 80%, and 90% sparsity settings. We report latency results for prefill (with 1K and 4K input length) as well as for generation with a single input token. Latency results are normalized to the FlashAttention baseline runtime for prefill (and to our Triton FlashDecoding baseline for generation) for the same input length.

We used PG-19 data since language modeling data allowed us to segment the fixed context and input into the desired lengths for measuring latency. Using real data ensured that we had a realistic sparsity distribution across different heads and layers. We ran clustering offline using offloading to collect data with a context length of 512K, and loaded this in one layer at a time in order to collect measurements. We report the average runtime across all layers in all our experimental results.

For prefill, we benchmarked our Triton kernel implementations using `triton.testing.do_bench` with 100 warmup runs and 500 measurement runs. For generation, we used 50 measurement runs for 100 different input query tokens, and averaged the runtime across all of these runs. We use the H100 NVL hardware platform for our experiments. We benchmarked the end-to-end runtime for our centroid lookup and sparse FlashAttention kernels (and include the runtime of Pytorch code for setting up arguments for our kernels). For prefill, we compared the performance of our implementation with the FlashAttention-2 implementation provided through the PyTorch `scaled_dot_product_attention` API. For generation, we implemented a Triton FlashDecoding kernel optimized for single-batch inference to serve as a stronger baseline.

6.3.2 Results

Figure 5 shows the latency for the FlashAttention baseline as well as for SQUEEZED ATTENTION with 70%, 80%, and 90% sparsity with 512K context length. We set the number of centroids to be 5% of the context length. We report results for generation (one input token) as well as prefill with 1K and 4K input tokens, and we normalize the latency to the baseline latency for each input size. These results show the benefits of our method for accelerating long context length inference, with $4.3\times / 4.2\times$ speedups demonstrated for the prefill and decode phases. This shows how our approach is flexible and can accelerate both prefill and generation stages. Additional results for context length 128K are provided in Appendix F.

7 Conclusion

In this work, we propose SQUEEZED ATTENTION as a method for accelerating attention in long context-length applications. Our method groups the keys and uses representative centroids for each group to quickly identify which keys are important for the attention operation. For fixed context applications, we can cluster the keys offline using K-means. Online during inference, we first compare the new input query with the representative centroids, and then only compute exact attention for these important keys. Our method can be extended to the hierarchical retrieval scheme, which can reduce the memory and compute complexity of lookups to logarithmic complexity with respect to the fixed context length. SQUEEZED ATTENTION is able to provide $4.3\times / 4.2\times$ speedups during prefill and decode

phases for long context inference, while maintaining accuracy. Additionally, we outline how our algorithm can be extended using a hierarchical centroid lookup, allowing us to achieve the accuracy of fine-grained centroid lookups while maintaining the efficiency of coarse-grained centroids, thereby improving the scalability of our approach for longer context lengths. Our approach accelerates long context length LLM inference with fixed context applications while maintaining accuracy.

Acknowledgements

We are grateful for the insightful discussions with Dhairya Malhotra. We acknowledge gracious support from the FuriosaAI team including Jihoon Yoon, Suyeol Lee, and Hyung Il Koo, as well as from Intel, Apple, and NVIDIA. We also appreciate the support from Microsoft through their Accelerating Foundation Model Research, including great support from Sean Kuno. Furthermore, we appreciate support from Google Cloud, the Google TRC team, and specifically Jonathan Caton, and Prof. David Patterson. Prof. Keutzer’s lab is sponsored by the Intel corporation, Intel One-API, Intel VLAB team, the Intel One-API center of excellence, as well as funding through BDD and BAIR. We appreciate great feedback and support from Ellick Chan, Saurabh Tangri, Andres Rodriguez, and Kittur Ganesh. Sehoon Kim would like to acknowledge the support from the Korea Foundation for Advanced Studies (KFAS). Michael W. Mahoney would also like to acknowledge a J. P. Morgan Chase Faculty Research Award as well as the DOE, NSF, and ONR. This work was supported by the Director, Office of Science, Office of Advanced Scientific Computing Research, of the U.S. Department of Energy under Contract No. DE-AC02-05CH11231. Our conclusions do not necessarily reflect the position or the policy of our sponsors, and no official endorsement should be inferred.

References

- [1] arxiv api user’s manual. <https://info.arxiv.org/help/api/index.html>. Accessed: 2023-10-13.
- [2] togethercomputer/llama-2-7b-32k, 2023.
- [3] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [4] Anthropic. Claude 2: <https://www.anthropic.com/news/claude-2>, 2023.
- [5] Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, et al. Longbench: A bilingual, multitask benchmark for long context understanding. *arXiv preprint arXiv:2308.14508*, 2023.
- [6] William Brandon, Mayank Mishra, Aniruddha Nrusimha, Rameswar Panda, and Jonathan Ragan Kelly. Reducing transformer key-value cache size with cross-layer attention. *arXiv preprint arXiv:2405.12981*, 2024.
- [7] Yukang Chen, Shengju Qian, Haotian Tang, Xin Lai, Zhijian Liu, Song Han, and Jiaya Jia. Longlora: Efficient fine-tuning of long-context large language models. *arXiv preprint arXiv:2309.12307*, 2023.
- [8] Rewon Child, Scott Gray, Alec Radford, and Ilya Sutskever. Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*, 2019.
- [9] Ronald Coifman, Vladimir Rokhlin, and Stephen Wandzura. The fast multipole method for the wave equation: A pedestrian prescription. *IEEE Antennas and Propagation magazine*, 35(3):7–12, 1993.
- [10] Tri Dao. Flashattention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*, 2023.
- [11] Tri Dao, Daniel Haziza, Francisco Massa, and Grigory Sisov. Flash-decoding for long-context inference: <https://crfm.stanford.edu/2023/10/12/flashdecoding.html>, 2023.

- [12] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- [13] Qichen Fu, Minsik Cho, Thomas Merth, Sachin Mehta, Mohammad Rastegari, and Mahyar Najibi. Lazyllm: Dynamic token pruning for efficient long context llm inference. *arXiv preprint arXiv:2407.14057*, 2024.
- [14] Suyu Ge, Yunan Zhang, Liyuan Liu, Minjia Zhang, Jiawei Han, and Jianfeng Gao. Model tells you what to discard: Adaptive kv cache compression for llms. *arXiv preprint arXiv:2310.01801*, 2023.
- [15] Google. Gemini 1.5 <https://blog.google/technology/ai/google-gemini-next-generation-model-february-2024>, 2023.
- [16] Alexander Gray and Andrew Moore. N-body problems in statistical learning. *Advances in neural information processing systems*, 13, 2000.
- [17] Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. Kvquant: Towards 10 million context length llm inference with kv cache quantization. *Advances in Neural Information Processing Systems*, 2024.
- [18] Cheng-Ping Hsieh, Simeng Sun, Samuel Krman, Shantanu Acharya, Dima Rekesh, Fei Jia, and Boris Ginsburg. Ruler: What’s the real context size of your long-context language models? *arXiv preprint arXiv:2404.06654*, 2024.
- [19] Hao Kang, Qingru Zhang, Souvik Kundu, Geonhwa Jeong, Zaoxing Liu, Tushar Krishna, and Tuo Zhao. Gear: An efficient kv cache compression recipe for near-lossless generative inference of llm. *arXiv preprint arXiv:2403.05527*, 2024.
- [20] Yanming Kang, Giang Tran, and Hans De Sterck. Fast multipole attention: A divide-and-conquer attention mechanism for long sequences. *arXiv preprint arXiv:2310.11960*, 2023.
- [21] Minsoo Kim, Kyuhong Shim, Jungwook Choi, and Simyung Chang. Infinipot: Infinite context processing on memory-constrained llms. *arXiv preprint arXiv:2410.01518*, 2024.
- [22] Dongryeol Lee, Andrew Moore, and Alexander Gray. Dual-tree fast gauss transforms. *Advances in Neural Information Processing Systems*, 18, 2005.
- [23] Dacheng Li, Rulin Shao Shao, Anze Xie, Ying Sheng, Lianmin Zheng, Joseph Gonzalez, Ion Stoica, Xuezhe Ma, and Hao Zhang. How long can open-source llms truly promise on context length?, June 2023.
- [24] Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. Snapkv: Llm knows what you are looking for before generation. *arXiv preprint arXiv:2404.14469*, 2024.
- [25] Di Liu, Meng Chen, Baotong Lu, Huiqiang Jiang, Zhenhua Han, Qianxi Zhang, Qi Chen, Chengruidong Zhang, Bailu Ding, Kai Zhang, et al. Retrievalattention: Accelerating long-context llm inference via vector retrieval. *arXiv preprint arXiv:2409.10516*, 2024.
- [26] Hao Liu, Wilson Yan, Matei Zaharia, and Pieter Abbeel. World model on million-length video and language with ringattention. *arXiv preprint arXiv:2402.08268*, 2024.
- [27] Ruikang Liu, Haoli Bai, Haokun Lin, Yuening Li, Han Gao, Zhengzhuo Xu, Lu Hou, Jun Yao, and Chun Yuan. Intactkv: Improving large language model quantization by keeping pivot tokens intact. *arXiv preprint arXiv:2403.01241*, 2024.
- [28] Zichang Liu, Aditya Desai, Fangshuo Liao, Weitao Wang, Victor Xie, Zhaozhuo Xu, Anastasios Kyrillidis, and Anshumali Shrivastava. Scissorhands: Exploiting the persistence of importance hypothesis for llm kv cache compression at test time. *Advances in Neural Information Processing Systems*, 36, 2024.

- [29] Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. Kivi: A tuning-free asymmetric 2bit quantization for kv cache. *arXiv preprint arXiv:2402.02750*, 2024.
- [30] William B March, Bo Xiao, and George Biros. Askit: Approximate skeletonization kernel-independent treecode in high dimensions. *SIAM Journal on Scientific Computing*, 37(2):A1089–A1110, 2015.
- [31] Vlad Morariu, Balaji Srinivasan, Vikas C Raykar, Ramani Duraiswami, and Larry S Davis. Automatic online tuning for fast gaussian summation. *Advances in neural information processing systems*, 21, 2008.
- [32] Piotr Nawrot, Adrian Łańcucki, Marcin Chochowski, David Tarjan, and Edoardo M Ponti. Dynamic memory compression: Retrofitting llms for accelerated inference. *arXiv preprint arXiv:2403.09636*, 2024.
- [33] OpenAI. Introducing triton: Open-source gpu programming for neural networks: <https://openai.com/index/triton/>, 2021.
- [34] OpenAI. <https://github.com/triton-lang/triton/blob/main/python/tutorials/06-fused-attention.py>, 2024.
- [35] Matanel Oren, Michael Hassid, Yossi Adi, and Roy Schwartz. Transformers are multi-state rnns. *arXiv preprint arXiv:2401.06104*, 2024.
- [36] Matteo Pagliardini, Daniele Paliotta, Martin Jaggi, and François Fleuret. Fast attention over long sequences with dynamic sparse flash attention. *Advances in Neural Information Processing Systems*, 36:59808–59831, 2023.
- [37] Jack W Rae, Anna Potapenko, Siddhant M Jayakumar, Chloe Hillier, and Timothy P Lillicrap. Compressive transformers for long-range sequence modelling. arxiv preprint, 2019. URL <https://arxiv.org/abs/1911>.
- [38] Aurko Roy, Mohammad Saffar, Ashish Vaswani, and David Grangier. Efficient content-based sparse attention with routing transformers. *Transactions of the Association for Computational Linguistics*, 9:53–68, 2021.
- [39] Jiaming Tang, Yilong Zhao, Kan Zhu, Guangxuan Xiao, Baris Kasikci, and Song Han. Quest: Query-aware sparsity for efficient long-context llm inference. *arXiv preprint arXiv:2406.10774*, 2024.
- [40] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. LLaMA: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- [41] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- [42] Yang, Duraiswami, and Gumerov. Improved fast gauss transform and efficient kernel density estimation. In *Proceedings ninth IEEE international conference on computer vision*, pages 664–671. IEEE, 2003.
- [43] Yao Yao, Zuchao Li, and Hai Zhao. Sirllm: Streaming infinite retentive llm. *arXiv preprint arXiv:2405.12528*, 2024.
- [44] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, et al. Big bird: Transformers for longer sequences. *Advances in neural information processing systems*, 33:17283–17297, 2020.
- [45] Hailin Zhang, Xiaodong Ji, Yilin Chen, Fangcheng Fu, Xupeng Miao, Xiaonan Nie, Weipeng Chen, and Bin Cui. Pqcache: Product quantization-based kvcache for long context llm inference. *arXiv preprint arXiv:2407.12820*, 2024.
- [46] Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, et al. H2o: Heavy-hitter oracle for efficient generative inference of large language models. *Advances in Neural Information Processing Systems*, 36, 2024.

A Additional Analysis

A.1 t-SNE Visualization of Keys and Their Clusters




Figure A.1: t-SNE visualization of key embeddings and their Level 1 and 2 clusters from LLaMA-2-7B-32K on the TREC benchmark (two attention heads, with index 24 and 25, from layer 0). For clarity, only the top 15 Level 1 clusters nearest to the query are shown.

Figure A.1 illustrates t-SNE plots of key embeddings and their Level 1 and 2 clusters. As can be seen, while the coarser Level 2 clusters offer a rough grouping of the keys, the finer Level 1 clusters allow for a more detailed and accurate representation within each cluster.

A.2 Attention Score Skewness Analysis




Figure A.2: Cumulative attention scores for the top 1% highest scoring attention values across different heads and layers in LLaMA-2-7B-32K for a single sample on the TREC benchmark.

Figure A.2 illustrates the cumulative attention scores for the top 1% highest attention values within the same model. A higher value (up to a maximum of 1) indicates that the attention head has a sharper, more skewed distribution, while

a lower value indicates a flatter distribution of attention scores. This plot demonstrates how the attention heads in the first two layers of the LLaMA-2-7B-32K model, as well as a subset of the heads at each of the remaining layers, have a flatter distribution of attention scores, and therefore we need to load in more keys for accurate computation of attention.

B Generation Centroid Lookup Kernel Ablation

We provide ablations for our centroid lookup kernel implementation for generation. Specifically, we ablate the benefits of our single-pass optimization, as well as the improvements from parallelizing along the KV sequence length dimension in order to accelerate the centroid lookup during generation. Table B.1 shows the results for this ablation, demonstrating how these optimizations allow SQUEEZED ATTENTION to achieve greater speedups during generation.

Table B.1: Ablation for our centroid lookup kernel implementation during generation with sequence length 512K, showing the normalized latency relative to the baseline Triton FlashDecoding kernel. We show the benefits of incorporating our single-pass optimization, as well as the gains from splitting along the KV dimension as in FlashDecoding [11]. These results highlight the importance of our lookup kernel optimizations for attaining speedups with SQUEEZED ATTENTION during generation.

Configuration	FlashDecoding Baseline	Centroid Lookup	+ Single-Pass Optimization	+ Split-KV Optimization
Normalized Latency	1	0.29	0.22	0.12

C Experimental Details

KV Budget Computation. We report KV budget estimates throughout the evaluation based on the configured sparsity threshold and percentage of centroids used. Note that the KV budget does not include performing recomputation with the same key centroid (as our current kernel implementation for prefill loads the key centroid twice to avoid materializing intermediate tensors). Additionally, due to our calibration procedure (which sets a single threshold for both prefill and generation), the KV cache budget may be slightly higher than expected during prefill, and slightly lower than expected during generation. This occurs since averaging the attention to the centroids across query tokens flattens the attention distribution, which leads to preserving more key tokens during prefill. Also, with the hierarchical method, the portion of KV tokens that are loaded may deviate further from the expected value from calibration due to the potential for incorrectly filtering out important keys when comparing with the Level 1 centroids as well as due to not loading all of the Level 2 keys when computing the denominator in Equation 3. For our hierarchical lookup experiments, we therefore profiled the KV cache budget estimates reported in our evaluation.

Baseline Methods. QUEST [39] uses a fixed token budget across all input samples when evaluating on LongBench. In order to perform a fair comparison with our method, we set the token budget for their approach to be a fixed percentage of the length of each sample. We dynamically set their token budget to be a percentage of the fixed context length (rounded up to the nearest multiple of 16). Since our method only approximates attention to the fixed-context portion of the input, we adjust this dynamically computed token budget using the full input length for the sample as well as the maximum generation length for the target task. This adjustment ensures that the achieved compression ratio for the fixed context with their method is comparable with our approach. Note that this adjustment also accounts for the 100 tokens at the end of the fixed context that we use for calibration purposes and therefore retain exactly with our method. For QUEST comparisons, we also leave the first two layers uncompressed to match their default configuration. We use 90% and 95% sparsity settings for their method to obtain the two configurations reported in the paper.

Implementation Details. Similar to the numerically stable implementation of Softmax, where the maximum value is subtracted from all inputs, our centroid lookup approach during the generation stage (Section 4.1) also subtracts the maximum value from all inputs while computing the denominator D . Then, when comparing $\exp(qC_i^\top)$ to the threshold DT , the maximum value correction can similarly be accounted for by scaling the threshold DT using the exponential of the maximum value rather than by scaling $\exp(qC_i^\top)$ down by this value.

D Data Curation Prompts

Below we provide the GPT-4-Turbo [3] prompts that were used to generate and filter high-quality question-answer pairs for our custom benchmark PreFixQA.

D.1 Dataset Curation Prompts for Question Generation

System: You are a helpful assistant tasked with generating very specific, unambiguous, short-answer questions based on a provided Document. The goal is to create high-quality synthetic data. Ensure the following:

1. The question must be fact-based and directly answerable from the text of the provided Document and section of the paper.
2. The question should not be vague, subjective, or open to interpretation.
3. Focus on creating concise, precise questions that yield a 1-2 word answer.
4. Questions should be relevant to the section of the paper and avoid overly broad phrasing.
5. Avoid generating questions where the answer is too complex or requires long explanations.

User: Based on the section of the paper from the given document which is an arxiv paper generate one short-answer question that asks for specific information retrievable directly from the section of the paper. The answer must be 1-2 words only.

Follow the format of this example:

Example: Question: What type of buffer is integrated with warp-specialization in FlashAttention-3? Use this section of the paper as the context: "Given section". Do not output the answer; Just the question.

D.2 Dataset Curation Prompts for Answer Generation

System: You are a helpful assistant designed to generate accurate and specific short answers from a provided section of the paper and Document. Ensure the following:

1. The answer must be concise (1-2 words).
2. The answer must be directly retrieved from the provided text.
3. If the section of the paper does not contain the information necessary to answer the question, respond with: 'The document does not contain the answer to the question.'.
4. Avoid providing additional commentary, and only output the answer.

User: Given the section of the paper below from an arXiv paper, generate a concise (1-2 words) answer to the following question. Retrieve the answer from the paper and the provided paragraph. Pay close attention to the document and retrieve the right answer. Output the answer only and not the question or anything else.

Follow the format of this example:

Example: Question: What type of buffer is integrated with warp-specialization in FlashAttention-3? Answer: circular SMEM buffer. Here is the section of the paper: "Given Section". Question: "Question". Answer:

D.3 Dataset Curation Prompts - Filtering

System: Please act as an impartial judge and evaluate the quality of the question and answer pairs provided. You will be assessing them based on the provided document. Your evaluation should emphasize the following criteria:

1. Correctness: Is the answer factually accurate and grounded in the document?
2. Agreement: Does the answer directly address the question and provide a relevant response?
3. Confidence: Does the answer confidently engage with the question, even if the Document does not contain the exact information?

Important considerations for rating:

- Rate vague or overly general questions lower, especially if they lack specificity or do not make sense in the context of the document.
- Rate answers where the model, dataset, or method is unclear or missing details lower.
- If the answer states that the information is not in the document, confirm by reviewing the document. If the information is indeed missing, rate the answer highly. If it is present, rate the answer lower.
- Avoid focusing on questions about appendix numbers, or formatting details (like section names).
- Avoid asking questions that have 2 possible answers. If there are 2 possible answers and only one is provided, rate the answers low.
- If there are questions that are taken from the 'References' and 'Acknowledgments' sections, rate the answers low.

For each answer, provide a rating on a scale from 1 to 10 based on its quality.

User: The question is "Question" and the answers are "Answers". Make sure to give the rating for each answer in the answers list. Please output your rating in the following format:

Question: "Question" Answer: [Answer1]- Rating: [[5]] Answer: [Answer2]- Rating: [[8]] ...

E Comparison with Ideal Lookup

Table E.1 provides comparisons with a baseline using an idealized lookup. For the "Ideal" baseline comparisons, we first compute attention from the user input query tokens to all of the fixed context keys. We then select the keys whose attention scores are above the configured threshold, and compute exact attention using only these keys. This serves as an upper bound on the attainable accuracy for a given sparsity percentage, since it lets us identify the high-scoring tokens exactly before computing attention using these tokens. For this idealized baseline, we also calibrate for a global threshold across all layers to allow this approach to adaptively retain more or fewer keys for different heads, which allows us to make a fair comparison between our method and this idealized baseline.

Table E.1: Ablation showing the accuracy of SQUEEZED ATTENTION compared with an idealized baseline. We report LongBench evaluation results for the Llama-2-7B-32K model, including the average score across tasks. We also include the KV budget for our method (“Budget”), which gives the percentage of the KV cache that needs to be loaded in during inference (including extra memory movement for cluster centroids). Our results demonstrate that our centroid lookup method attains similar accuracy to the idealized baseline for the same level of sparsity across different LongBench tasks.

Config	Budget	Single-Document QA			Multi-Document QA			Summarization			Few-shot Learning			Code		Avg.
		NQA	Qasper	MFQA	Hotpot	2Wiki	Musique	GRp	QMSum	MNews	TREC	TQA	SSum	RBench	LCC	
All KV	1	17.91	11.12	33.87	12.45	11.95	6.54	29.37	16.93	21.58	71.50	87.96	43.87	61.45	59.14	34.69
Ideal-70%	0.3	18.28	11.25	34.13	12.56	12.13	6.57	29.12	16.99	21.37	70.00	87.79	43.59	62.01	59.36	34.65
Squeeze-70%	0.325	18.55	11.78	34.33	12.31	12.31	6.26	29.50	16.90	20.76	69.00	87.96	43.90	61.29	59.53	34.60
Ideal-90%	0.1	17.65	11.77	33.89	12.67	11.86	6.04	29.02	16.79	23.22	69.00	87.23	44.33	57.82	60.17	34.39
Squeeze-90%	0.125	18.15	14.39	32.38	11.84	11.70	6.45	29.06	16.93	21.66	70.00	87.43	45.15	58.79	59.37	34.52

F Kernel Benchmarking for 128K Sequence Length

Figure F.1 shows the latency for the FlashAttention baseline as well as for SQUEEZED ATTENTION with 70%, 80%, and 90% sparsity with 128k context length, with the number of centroids set to be 5% of the context length. We report results for generation as well as prefill with 1K and 4K input tokens, and we report the latency for each configuration normalized to the baseline latency for the corresponding input size. For prefill, we observe $4.2\times$ speedups with 90% sparsity, which is comparable with our speedups reported for 512K. For generation, we observe reduced speedups, observing $2.5\times$ speedups with 90% sparsity, relative to the FlashDecoding baseline implementation. The reduced speedups in this regime are due to greater overheads with the centroid lookup kernel for shorter sequence lengths.




Figure F.1: Kernel implementation latency results for FlashAttention baseline as well as for SQUEEZED ATTENTION with 70%, 80%, and 90% sparsity settings. We report latency results for prefill (with 1K and 4K input length) as well as for generation with a single input token. Latency results are normalized to the FlashAttention baseline runtime for prefill (and to our Triton FlashDecoding baseline for generation) for the same input length.

G Limitations

One of the limitations of our work is that both the sparsity threshold and the number of centroids used are hyperparameters. Furthermore, the degree of sparsity that is attainable without accuracy degradation is also dependent on the input context and the type of task. Our work could therefore be extended by developing an automated way of configuring

these hyperparameters depending on the target accuracy level and the input context. Our approach also focuses on accelerating fixed context applications, which limits its use for applications where the full context is only available online. Future work can be done to accelerate the initial offline clustering step in order to allow our method to be used in online use-cases.